

Empirical reliability and security of drift-aware continuous training

Companion to JSSOFTWARE-D-26-02282 · kind dact-local-eks + Compose window · measured

Reliability (kind, n=5)

Backend pod: 1/5 /ready in 30 s

MLflow down: train hung ≥ 45 s

MinIO not deployed on kind

Rollback v7→v1 in 0.76 s

Security (scan + STRIDE)

86 OSV advisories / 18 pins

10 credential / port findings

5 of 7 High STRIDE boundaries

Silent MLflow fail confirmed live

A/B/C (kind, 9 runs)

C fired on feature_drift

C held on healthy / regression

B waits avg 3.5 scheduled days

Compose: 0/9 pass; F1-only 6/9